

# Clues Not Processing Troubleshooting Guide

Monday, 24 January 2022 11:18 AM

24 Jan 2022 uploaded to <https://github.com/CluedIn-io/ImpTools/tree/master/k8s/pwsh>

## Purpose

To ensure clue processing is flowing through the system.

## Things to Check

Methods:

- Connect to your k8s cluster with Lens (extremely helpful! - see <https://k8slens.dev/>)
- Use kubectl on the command line with powershell - see <https://github.com/CluedIn-io/ImpTools/tree/master/k8s/pwsh> for some sample scripts
- Azure Portal (<https://portal.azure.com/>)


Check the Overall status

RabbitMQ (<https://www.rabbitmq.com/>) pod coming up before the volumn is ready - restart the pod


Warnings: 1			
Message	Object ^	Type ^	Age ^
Unable to attach or mount volumes: u...	cluedin-demo-rabbitmq-0	Pod	103m

Pod is waiting for another pod to be available - this is how the start up dependancies are expressed.

In this case go sort out the gql pod first.



If all is okay pod wise, expect to see this



Also in the Azure Portal - check the workloads

APAC-demo-k8s | Workloads

Kubernetes service

Search (Ctrl+F)

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Security

Kubernetes resources

Namespaces Workloads Services and ingresses Storage Configuration Settings Node pools Cluster configuration Networking GitOps (preview) Deployment center (preview) Policies Properties Locks Monitoring Insights Alerts

Deployments Pods Replica sets Stateful sets Daemon sets Jobs Cron jobs

Filter by deployment name Filter by label selector Filter by namespace

Enter the full deployment name foo=bar,key=value cluedin

Name	Namespace	Ready	Up-to-date	Available	Age
haproxy-ingress	cluedin	1/1	1	1	48 days
cluedin-demo-annotation	cluedin	1/1	1	1	47 days
cluedin-demo-clean	cluedin	1/1	1	1	47 days
cluedin-demo-cluedin	cluedin	1/1	1	1	47 days
cluedin-demo-cluedincontroller	cluedin	1/1	1	1	47 days
cluedin-demo-datasource	cluedin	1/1	1	1	47 days
cluedin-demo-gql	cluedin	1/1	1	1	47 days
cluedin-demo-grafana	cluedin	1/1	1	1	47 days
cluedin-demo-kube-state-metrics	cluedin	1/1	1	1	47 days
cluedin-demo-monitoring-operator	cluedin	1/1	1	1	47 days
cluedin-demo-neo4j	cluedin	1/1	1	1	47 days
cluedin-demo-openrefine	cluedin	1/1	1	1	47 days
cluedin-demo-processing	cluedin	1/1	1	1	47 days
cluedin-demo-redis	cluedin	1/1	1	1	47 days
cluedin-demo-sqlserver	cluedin	1/1	1	1	47 days
cluedin-demo-submitter	cluedin	1/1	1	1	47 days
cluedin-demo-ui	cluedin	1/1	1	1	47 days
cluedin-demo-webapi	cluedin	1/1	1	1	47 days

Using kubectl

<https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/gp.ps1> - get pods

check for restarts as a sign of issues and also the STATUS column  
(in this case we had some issues, but they have been resolved).

```
PS C:\rudi.harris\src\CluedIn-io\ImpTools\k8s\pwsh> cat .\gp.ps1
kubectl get pods
PS C:\rudi.harris\src\CluedIn-io\ImpTools\k8s\pwsh> .\gp.ps1
NAME                                READY   STATUS    RESTARTS   AGE
alertmanager-cluedin-demo-monitoring-alertmanager-0  2/2     Running   0           2d21h
cluedin-demo-annotation-84764c8df9-6tnvk             1/1     Running   0           2d21h
cluedin-demo-clean-66b64c7578-b8z4w                 1/1     Running   0           2d21h
cluedin-demo-cluedin-9f6cf9c6d-r8z5                  1/1     Running   3           2d21h
cluedin-demo-cluedincontroller-65bd75f4b9-9sdzd       1/1     Running   0           2d21h
cluedin-demo-datasource-6b75cd87d8-mb977            1/1     Running   0           2d21h
cluedin-demo-gql-5bf88c9dd9-5jbvt                    1/1     Running   0           2d21h
cluedin-demo-grafana-86c89f8454-gn82s                2/2     Running   0           2d21h
cluedin-demo-kube-state-metrics-5d5bdd975-w2pn6      1/1     Running   0           2d21h
cluedin-demo-monitoring-operator-555c9bff9c-pm49b    1/1     Running   0           2d21h
cluedin-demo-neo4j-5c89778749-8hqdg                 1/1     Running   0           2d21h
cluedin-demo-openrefine-58d66f5979-g2qmr             1/1     Running   0           2d21h
cluedin-demo-processing-764659974b-r8bqx             1/1     Running   1           2d21h
cluedin-demo-prometheus-node-exporter-f5vz4          1/1     Running   0           3h30m
cluedin-demo-prometheus-node-exporter-ndcmf          1/1     Running   0           3h29m
cluedin-demo-prometheus-node-exporter-p6r7l          1/1     Running   0           3h30m
cluedin-demo-prometheus-node-exporter-sx658          1/1     Running   0           3h29m
cluedin-demo-prometheus-node-exporter-wd855          1/1     Running   0           3h29m
cluedin-demo-prometheus-node-exporter-x2fms          1/1     Running   0           3h30m
cluedin-demo-rabbitmq-0                              1/1     Running   0           8m51s
cluedin-demo-redis-78686bc979-jmkkz                 1/1     Running   0           2d21h
cluedin-demo-sqlserver-847f8c6594-2b767             1/1     Running   0           2d21h
cluedin-demo-submitter-5bbcd54774-7tqsv             1/1     Running   0           2d21h
cluedin-demo-ui-6fbfdcbf7-5gcr6                     1/1     Running   0           2d21h
cluedin-demo-webapi-5686697599-gwqcp                 1/1     Running   0           2d21h
elasticsearch-master-0                              1/1     Running   0           2d21h
haproxy-ingress-5c9cb6796f-p9v6n                    1/1     Running   0           2d21h
prometheus-cluedin-demo-monitoring-prometheus-0      2/2     Running   1           2d21h
```

Check logs

Pod: cluedin-demo-neo4j-5c89778749-8hqdg

Pod Logs

CPU Memory Network Filesystem

0.015

Main server pod, processing pod, neo4j pod - these are the best 3 to start with  
Look for out of memory issues in neo4j  
Look for C# error stacks and errors in the others

```

Pod cluedin-demo-cluedin-9f6cf9c6d-r... X Pod cluedin-demo-processing-764659... X Pod cluedin-demo-neo4j-5c89778749-... X +
Namespace cluedin Pod cluedin-demo-neo4j Container neo4j Search...
2022-01-24 01:26:37.919+0000 WARN Unknown config option: causal_clustering.transaction_listen_address
2022-01-24 01:26:37.928+0000 INFO ===== Neo4j 3.5.30 =====
2022-01-24 01:26:37.933+0000 INFO Starting...
Creating initial indexes and constraints
Connection refused
2022-01-24 01:26:44.918+0000 INFO Bolt enabled on 0.0.0.0:7687.
2022-01-24 01:26:46.034+0000 INFO Started.
2022-01-24 01:26:46.152+0000 WARN Acceptors should be <= availableProcessors: ServerConnector@7b74883{HTTP/1.1, (http/1.1){0.0.0.0:0}}
2022-01-24 01:26:46.160+0000 WARN Acceptors should be <= availableProcessors: ServerConnector@280b9922{SSL, (ssl, http/1.1){0.0.0.0:0}}
2022-01-24 01:26:46.793+0000 INFO Remote interface available at http://localhost:7474/
Creating initial indexes and constraints
0 rows available after 25 ms, consumed after another 0 ms

```

Alternatively, use the kubectl command under git bash (so the coloured logs are shown correctly - <https://gitforwindows.org/>)

-l can be used to reference the pod using a label - or use the pod name from the get pods command

The below example is show the logs for the server (i.e. role=main) starting from 15 minutes ago and -f for follow the log as new entries come in

```

MINGW64/c/Users/rudi
rudi@Rudiwin10MBP MINGW64 ~
$ kubectl logs -l role=main --since=15m -f
[05:05:10 VRB][CluedIn.Core.ExecutionContext] Job Queue : [0] - OrgId: [null] - AgentId: [null] - AgentGroupId:
[d1238808-b8d4-4fd3-a5f7-959f5a1895a2]
[05:05:10 VRB][CluedIn.Server.Agent.AgentControllerComponent] [0] AgentGroupJobs found.
[05:05:10 VRB][CluedIn.Core.ExecutionContext] Job Queue : [0] - OrgId: [null] - AgentId: [null] - AgentGroupId:
[null]
[05:05:10 VRB][CluedIn.Server.Agent.AgentControllerComponent] [0] SharedJobs found.
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDDV" accepted.
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDDV" started.
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel.Transport.Sockets] Connection id "0HMEUR51HUDDV" received FI
N.
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDDV" disconnecting.
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel.Transport.Sockets] Connection id "0HMEUR51HUDDV" sending FIN
because: "The client closed the connection."
[05:05:12 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDDV" stopped.
[05:05:14 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDE0" accepted.
[05:05:14 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDE0" started.
[05:05:14 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDE0" disconnecting.
[05:05:14 DBG][Microsoft.AspNetCore.Server.Kestrel] Connection id "0HMEUR51HUDE0" stopped.
[05:05:14 DBG][Microsoft.AspNetCore.Server.Kestrel.Transport.Sockets] Connection id "0HMEUR51HUDE0" sending FIN
because: "The Socket transport's send loop completed gracefully."
[05:05:15 VRB][CluedIn.Core.ExecutionContext] Job Queue : [0] - OrgId: [null] - AgentId: [cd71757d-17a1-4d85-92
b3-f525da611ae2] - AgentGroupId: [null]
[05:05:15 VRB][CluedIn.Core.ExecutionContext] Job Queue : [0] - OrgId: [null] - AgentId: [null] - AgentGroupId: [null]

```


Here is a few examples of labels:

```

rudi@Rudiwin10MBP MINGW64 /c/rudi.harris/src/CluedIn-io/ImpTools/k8s/pwsh (master)
$ grep "\-l" restart*.ps1
restart_ingress_ui.ps1:kubectl delete pod -l app=kubernetes.io/name=haproxy-ingress
restart_ingress_ui.ps1:kubectl delete pod -l app=ui
restart_main_processing_crawling.ps1:kubectl delete pod -l role=main
restart_main_processing_crawling.ps1:kubectl delete pod -l role=processing
restart_main_processing_crawling.ps1:kubectl delete pod -l role=crawling
restart_rmqs.ps1:kubectl delete pod -l app=kubernetes.io/name=rabbitmq
restart_submitter.ps1:kubectl delete pod -l app=submitter

```

These can be found under labels in lens



or use `kubectl describe pod <pod name>`

and search for the labels section:

```


Labels:
  app=cluedin
  pod-template-hash=9f6cf9c6d
  release=cluedin-demo
  role=main

```

Check the queues in RabbitMQ

Pods						
Filtered: 1 / 28			Namespace: cluedin		rabbit	
<input type="checkbox"/> Name	Containers	Restarts	Node	Age	Status	
<input type="checkbox"/> cluedin-demo-rabbitmq-0	<div></div>	0	aks-datapool-175905	29m	Running	

Select then scroll to the Ports section



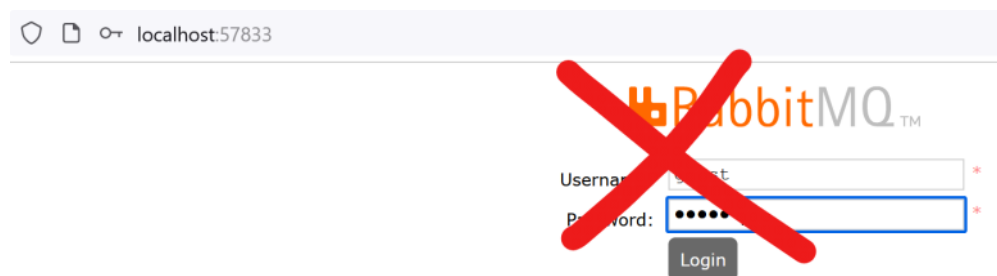
Click Forward for the stats port - put 15672 as the port then click Start

Or do the port forwarding using kubectl


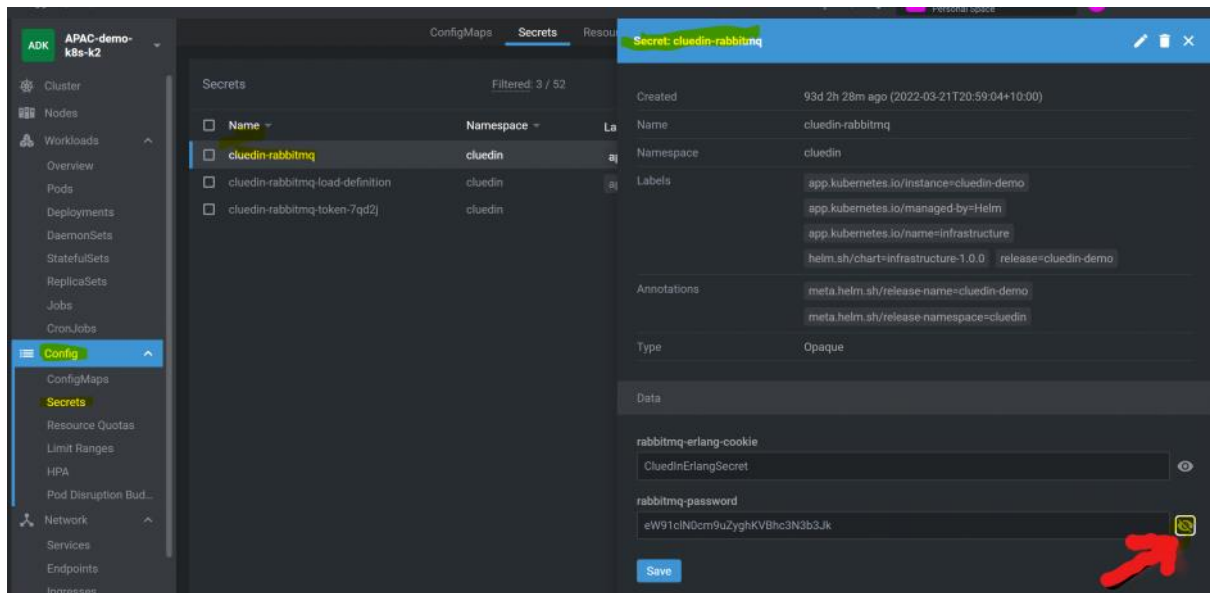
```
https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/pf_rmqs.ps1
PS C:\rudi.harris\src\CluedIn-io\ImpTools\k8s\pwsh> .\pf_rmqs.ps1
label is "app.kubernetes.io/name=rabbitmq", looking up name...
name is "pod/cluedin-demo-rabbitmq-0"
Forwarding from 127.0.0.1:15672 -> 15672
Forwarding from [::1]:15672 -> 15672
```

And open your browser to <http://localhost:15672/>

Login guest/guest




Update: from 3.3.0 onwards the username is cluedin and the password is a cluster secret "cluedin-rabbitmq"



Username:

Password:

Select queues



Check all the process command queues with the Regex filter for **.\*Processing.Process.\*Command**

Filter:  ☒ Regex

Overview				Messages			Message rates			
Name	Type	Features	State	Ready	Unacked	Total	incoming	deliver	/ get	ack
CluedIn.Core.Messages.Processing.ProcessBigClueCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				
CluedIn.Core.Messages.Processing.ProcessEdgesPriorityCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				
CluedIn.Core.Messages.Processing.ProcessLowPriorityClueCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				
CluedIn.Core.Messages.Processing.ProcessPrioritizedClueCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				
CluedIn.Core.Messages.Processing.ProcessVersionHistoryCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				
CluedIn.Core.Messages.Processing.ProcessWebHookClueCommand, CluedIn.Core_CluedIn	classic	D lazy-queue	idle	0	0	0				

And also for the clue queue

Regex filter **^clue\$**

Page  of 1 - Filter:  ☒ Regex ?

Overview				Messages			Message rates			
Name	Type	Features	State	Ready	Unacked	Total	incoming	deliver	/ get	ack
clue	classic	D lazy-queue	idle	0	0	0				

If Ready is all zeros then the submitter has not submitted any clues

If Ready is non-zero then we expect that number to go down over time - if it is not going down then it's time to restart the processing pod - see Things to Try below

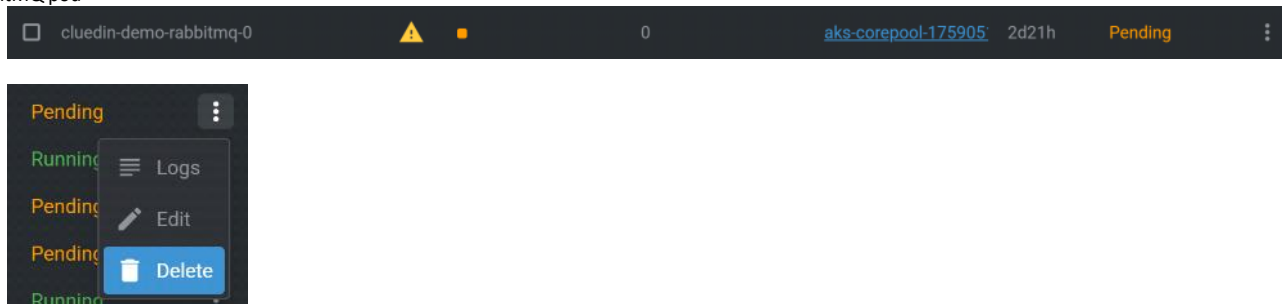
for the method

Restart processing pod

## Things to Try

Restarting pods

RabbitMQ pod



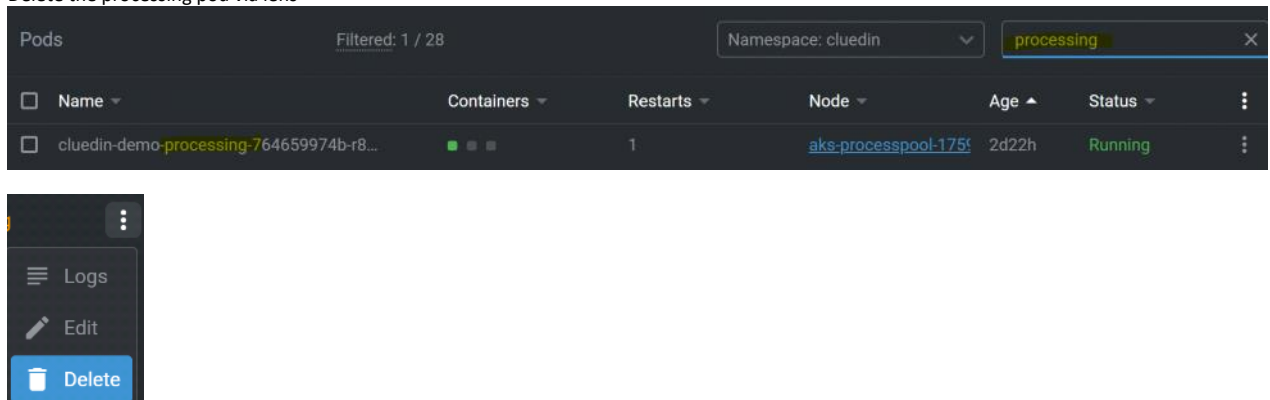
Or using kubectl and Powershell

[https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_rmqs.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_rmqs.ps1)

```
kubectl delete pod -l app.kubernetes.io/name=rabbitmq
```

Processing pod (and the other main pods)

Delete the processing pod via lens



Or using kubectl and Powershell

see the processing role line from [https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_main\\_processing\\_crawling.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_main_processing_crawling.ps1)

```
kubectl delete pod -l role=processing
```

Restart submitter pod

[https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_submitter.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_submitter.ps1)

```
kubectl delete pod -l app=submitter
```

Likewise, restart neo4j or gql

Restart by group

Restart the main pod (server), the processing pod and the crawling pod

[https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_main\\_processing\\_crawling.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_main_processing_crawling.ps1)

```
kubectl delete pod -l role=main
kubectl delete pod -l role=processing
kubectl delete pod -l role=crawling
```

Restart the Ingress and UI

[https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_ingress\\_ui.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_ingress_ui.ps1)

## If All Else Fails

[https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart\\_main\\_processing\\_crawling.ps1](https://github.com/CluedIn-io/ImpTools/blob/master/k8s/pwsh/restart_main_processing_crawling.ps1)

```
kubectl delete pod -l role=main
kubectl delete pod -l role=processing
kubectl delete pod -l role=crawling
```

Likewise, restart neo4j or gql

Hit the Stop button on the whole cluster

The screenshot shows the Microsoft Azure portal interface. At the top, the browser address bar displays the URL <https://portal.azure.com/#@cluedin.com/resource/subscriptions/3fa2142d-0d79-4df1>. The Azure portal header is blue with the 'Microsoft Azure' logo and a search bar. Below the header, the breadcrumb 'Home >' is visible. The main content area shows the 'APAC-demo-k8s' Kubernetes service, with a yellow highlight under the 'Kubernetes service' label. A toolbar contains buttons for 'Connect', 'Start', 'Stop' (highlighted in yellow), 'Delete', 'Refresh', and 'Give feedback'. On the left, a sidebar lists 'Overview' (selected) and 'Activity log'. On the right, under the 'Essentials' section, it shows 'Resource group : [APAC-demo-k8s-rg](#)' and 'Status : Succeeded (Running)'.

Wait until you are notified that everything has stopped.

Wait some more...

Then when you think it's ready, hit the Start button and start from the top of this advice to troubleshoot the pods as they start up - this takes 15 to 20 minutes typically.